


In [39]:

```
import numpy as np
import matplotlib.pyplot as plt

## Data generation
# write your code here
a=np.random.multivariate_normal([0.5,0], [[1, 0],[0, 1]], 100)
b=np.random.multivariate_normal([5,5], [[1, 0],[0, 1]], 100)
c=np.random.multivariate_normal([5,1], [[1, 0],[0, 1]], 100)
d=np.random.multivariate_normal([10,1.5], [[1, 0],[0, 1]], 100)
X=np.concatenate((a,b,c,d), axis=0)
plt.scatter(X[:,0],X[:,1])

n=X.shape[1]
n_iter=10
```



In [40]:

```
K=4
import random

# creating an empty centroid array
centroids=np.array([]).reshape(n,0)

# creating 5 random centroids
for k in range(K):
    centroids=np.c_[centroids,X[random.randint(0,m-1)]]

print(centroids)

[[ 6.03389968  1.07922277  6.52521602 -0.5505451 ]
 [ 0.52556462  0.84542642  1.00397515  1.56554586]]
```

In [41]:

```
output={}

# creating an empty array
euclid=np.array([]).reshape(m,0)

# finding distance between for each centroid
for k in range(K):
    dist=np.sum((X-centroids[:,k])**2,axis=1)
    euclid=np.c_[euclid,dist]

# storing the minimum value we have computed
minimum=np.argmin(euclid,axis=1)+1
```

In [42]:

```
# computing the mean of separated clusters
cent={}
for k in range(K):
    cent[k+1]=np.array([]).reshape(2,0)

# assigning of clusters to points
for k in range(m):
    cent[minimum[k]]=np.c_[cent[minimum[k]],X[k]]
for k in range(K):
    cent[k+1]=cent[k+1].T


# computing mean and updating it
for k in range(K):
    centroids[:,k]=np.mean(cent[k+1],axis=0)
```


In [43]:

```
# repeating the above steps again and again
for i in range(n_iter):
    euclid=np.array([]).reshape(m,0)
    for k in range(K):
        dist=np.sum((X-centroids[:,k])**2,axis=1)
        euclid=np.c_[euclid,dist]
    C=np.argmin(euclid,axis=1)+1
    cent={}
    for k in range(K):
        cent[k+1]=np.array([]).reshape(2,0)
    for k in range(m):
        cent[C[k]]=np.c_[cent[C[k]],X[k]]
    for k in range(K):
        cent[k+1]=cent[k+1].T
    for k in range(K):
        centroids[:,k]=np.mean(cent[k+1],axis=0)
    final=cent
```

In [44]:

```
for k in range(K):
    plt.scatter(final[k+1][:,0],final[k+1][:,1])
plt.scatter(centroids[0,:],centroids[1,:],s=300,c='yellow')
plt.rcParams.update({'figure.figsize':(10,7.5), 'figure.dpi':100})
plt.show()
```






In [45]:

```
import numpy as np
import matplotlib.pyplot as plt


## Data generation
# write your code here
a=np.random.multivariate_normal([0.5,0], [[1, 0],[0, 1]], 100)
b=np.random.multivariate_normal([5,5], [[1, 0],[0, 1]], 100)
c=np.random.multivariate_normal([5,1], [[1, 0],[0, 1]], 100)
d=np.random.multivariate_normal([10,1.5], [[1, 0],[0, 1]], 100)
X=np.concatenate((a,b,c,d), axis=0)
plt.scatter(X[:,0],X[:,1])

n=X.shape[1]
n_iter=10
```



In [47]:

```
from sklearn.mixture import GaussianMixture
gmm = GaussianMixture(n_components=4
                      ).fit(X)
labels = gmm.predict(X)
plt.scatter(X[:, 0], X[:, 1], c=labels, s=40, cmap='viridis');
probs = gmm.predict_proba(X)
```



In [48]:

```
from sklearn.cluster import KMeans
def calculate_mean_covariance(X, prediction):
    C = 3
    d = X.shape[1]
    labels = np.unique(prediction)
    initial_means = np.zeros((C, d))
    initial_cov = np.zeros((C, d, d))
    initial_pi = np.zeros(C)

    counter=0
    for label in sorted(labels):
        ids = np.where(prediction == label) # returns indices
        initial_pi[counter] = len(ids[0]) / X.shape[0]
        initial_means[counter,:] = np.mean(X[ids], axis = 0)
        de_meaned = X[ids] - initial_means[counter,:]
        Nk = X[ids].shape[0]
        initial_cov[counter,:, :] = np.dot(initial_pi[counter] * de_meaned.T, de_meaned)
    / Nk
    counter+=1
    assert np.sum(initial_pi) == 1
    return (initial_means, initial_cov, initial_pi)
```

```

n_clusters = 3
kmeans = KMeans(n_clusters= n_clusters, max_iter=500, algorithm = 'auto')
fitted = kmeans.fit(X)
prediction = kmeans.predict(X)

m, c, pi = calculate_mean_covariance(X, prediction)

```

In [49]:

```

from scipy.stats import multivariate_normal as mvn
class GMM:
    """ Gaussian Mixture Model

    Parameters
    -----
        k: int , number of gaussian distributions

        seed: int, will be randomly set if None

        max_iter: int, number of iterations to run algorithm, default: 200

    Attributes
    -----
        centroids: array, k, number_features

        cluster_labels: label for each data point

    """
    def __init__(self, C, n_runs):
        self.C = C # number of Guassians/clusters
        self.n_runs = n_runs

    def get_params(self):
        return (self.mu, self.pi, self.sigma)

    def calculate_mean_covariance(self, X, prediction):
        """Calculate means and covariance of different
            clusters from k-means prediction

        Parameters:
        -----
        prediction: cluster labels from k-means

        X: N*d numpy array data points

        Returns:
        -----
        intial_means: for E-step of EM algorithm

        intial_cov: for E-step of EM algorithm

    """
    d = X.shape[1]
    labels = np.unique(prediction)
    self.initial_means = np.zeros((self.C, d))
    self.initial_cov = np.zeros((self.C, d, d))
    self.initial_pi = np.zeros(self.C)

    counter=0
    for label in labels:
        ids = np.where(prediction == label) # returns indices
        self.initial_pi[counter] = len(ids[0]) / X.shape[0]
        self.initial_means[counter,:] = np.mean(X[ids], axis = 0)
        de_meaned = X[ids] - self.initial_means[counter,:]
        Nk = X[ids].shape[0] # number of data points in current gaussian
        self.initial_cov[counter,:, :] = np.dot(self.initial_pi[counter] * de_meaned
        .T, de_meaned) / Nk

```

```

        counter+=1
    assert np.sum(self.initial_pi) == 1

    return (self.initial_means, self.initial_cov, self.initial_pi)

def _initialise_parameters(self, X):
    """Implement k-means to find starting
        parameter values.
        https://datascience.stackexchange.com/questions/11487/how-do-i-obtain-the-weight-and-variance-of-a-k-means-cluster

    Parameters:
    -----
    X: numpy array of data points

    Returns:
    -----
    tuple containing initial means and covariance

    _initial_means: numpy array: (C*d)

    _initial_cov: numpy array: (C,d*d)

    """
    n_clusters = self.C
    kmeans = KMeans(n_clusters= n_clusters, init="k-means++", max_iter=500, algorithm = 'auto')
    fitted = kmeans.fit(X)
    prediction = kmeans.predict(X)
    self._initial_means, self._initial_cov, self._initial_pi = self.calculate_mean_covariance(X, prediction)

    return (self._initial_means, self._initial_cov, self._initial_pi)

def _e_step(self, X, pi, mu, sigma):
    """Performs E-step on GMM model

    Parameters:
    -----
    X: (N x d), data points, m: no of features
    pi: (C), weights of mixture components
    mu: (C x d), mixture component means
    sigma: (C x d x d), mixture component covariance matrices

    Returns:
    -----
    gamma: (N x C), probabilities of clusters for objects
    """
    N = X.shape[0]
    self.gamma = np.zeros((N, self.C))

    const_c = np.zeros(self.C)

    self.mu = self.mu if self._initial_means is None else self._initial_means
    self.pi = self.pi if self._initial_pi is None else self._initial_pi
    self.sigma = self.sigma if self._initial_cov is None else self._initial_cov

    for c in range(self.C):
        # Posterior Distribution using Bayes Rule
        self.gamma[:,c] = self.pi[c] * mvn.pdf(X, self.mu[c,:], self.sigma[c])

    # normalize across columns to make a valid probability
    gamma_norm = np.sum(self.gamma, axis=1)[:,np.newaxis]
    self.gamma /= gamma_norm

```

```

return self.gamma

def _m_step(self, X, gamma):
    """Performs M-step of the GMM
    We need to update our priors, our means
    and our covariance matrix.

    Parameters:
    -----
    X: (N x d), data
    gamma: (N x C), posterior distribution of lower bound

    Returns:
    -----
    pi: (C)
    mu: (C x d)
    sigma: (C x d x d)
    """
    N = X.shape[0] # number of objects
    C = self.gamma.shape[1] # number of clusters
    d = X.shape[1] # dimension of each object

    # responsibilities for each gaussian
    self.pi = np.mean(self.gamma, axis = 0)

    self.mu = np.dot(self.gamma.T, X) / np.sum(self.gamma, axis = 0)[: , np.newaxis]

    for c in range(C):
        x = X - self.mu[c, :] # (N x d)

        gamma_diag = np.diag(self.gamma[:,c])
        x_mu = np.matrix(x)
        gamma_diag = np.matrix(gamma_diag)

        sigma_c = x.T * gamma_diag * x
        self.sigma[c,:,:]=(sigma_c) / np.sum(self.gamma, axis = 0)[: , np.newaxis][c]

    return self.pi, self.mu, self.sigma

def _compute_loss_function(self, X, pi, mu, sigma):
    """Computes lower bound loss function

    Parameters:
    -----
    X: (N x d), data

    Returns:
    -----
    pi: (C)
    mu: (C x d)
    sigma: (C x d x d)
    """
    N = X.shape[0]
    C = self.gamma.shape[1]
    self.loss = np.zeros((N, C))

    for c in range(C):
        dist = mvn(self.mu[c], self.sigma[c], allow_singular=True)
        self.loss[:,c] = self.gamma[:,c] * (np.log(self.pi[c]+0.00001)+dist.logpdf(X
)-np.log(self.gamma[:,c]+0.00001))
    self.loss = np.sum(self.loss)
    return self.loss

def fit(self, X):
    """Compute the E-step and M-step and
    Calculates the lowerbound

    Parameters:

```

```

-----
X: (N x d), data

Returns:
-----
instance of GMM

"""

d = X.shape[1]
self.mu, self.sigma, self.pi = self._initialise_parameters(X)

try:
    for run in range(self.n_runs):
        self.gamma = self._e_step(X, self.mu, self.pi, self.sigma)
        self.pi, self.mu, self.sigma = self._m_step(X, self.gamma)
        loss = self._compute_loss_function(X, self.pi, self.mu, self.sigma)

        if run % 10 == 0:
            print("Iteration: %d Loss: %0.6f" %(run, loss))

except Exception as e:
    print(e)

return self

def predict(self, X):
    """Returns predicted labels using Bayes Rule to
    Calculate the posterior distribution

    Parameters:
    -----
    X: ?*d numpy array

    Returns:
    -----
    labels: predicted cluster based on
    highest responsibility gamma.

    """
    labels = np.zeros((X.shape[0], self.C))

    for c in range(self.C):
        labels[:,c] = self.pi[c] * mvn.pdf(X, self.mu[c,:], self.sigma[c])
    labels = labels.argmax(1)
    return labels

def predict_proba(self, X):
    """Returns predicted labels

    Parameters:
    -----
    X: N*d numpy array

    Returns:
    -----
    labels: predicted cluster based on
    highest responsibility gamma.

    """
    post_proba = np.zeros((X.shape[0], self.C))

    for c in range(self.C):
        # Posterior Distribution using Bayes Rule, try and vectorise
        post_proba[:,c] = self.pi[c] * mvn.pdf(X, self.mu[c,:], self.sigma[c])

    return post_proba

```


In [50]:

```
model = GMM(4, n_runs = 100)


fitted_values = model.fit(X)

predicted_values = model.predict(X)
# compute centers as point of highest density of distribution
centers = np.zeros((4,2))
for i in range(model.C):
    density = mvn(cov=model.sigma[i], mean=model.mu[i]).logpdf(X)
    centers[i, :] = X[np.argmax(density)]

plt.figure(figsize = (10,8))
plt.scatter(X[:, 0], X[:, 1],c=predicted_values ,s=50, cmap='viridis')

plt.scatter(centers[:, 0], centers[:, 1],c='black', s=300, alpha=0.6);
```

```
Iteration: 0 Loss: -1660.501398
Iteration: 10 Loss: -1653.847525
Iteration: 20 Loss: -1653.847330
Iteration: 30 Loss: -1653.847329
Iteration: 40 Loss: -1653.847329
Iteration: 50 Loss: -1653.847329
Iteration: 60 Loss: -1653.847329
Iteration: 70 Loss: -1653.847329
Iteration: 80 Loss: -1653.847329
Iteration: 90 Loss: -1653.847329
```



In [51]:

```
class FCM:
```

```

def __init__(self, n_clusters=10, max_iter=150, m=2, error=1e-5, random_state=42):
    assert m > 1
    self.u, self.centers = None, None
    self.n_clusters = n_clusters
    self.max_iter = max_iter
    self.m = m
    self.error = error
    self.random_state = random_state

def fit(self, X):
    self.n_samples = X.shape[0]
    r = np.random.RandomState(self.random_state)
    u = r.rand(self.n_samples, self.n_clusters)
    u = u / np.tile(u.sum(axis=1)[np.newaxis].T, self.n_clusters)

    r = np.random.RandomState(self.random_state)
    self.u = r.rand(self.n_samples, self.n_clusters)
    self.u = self.u / np.tile(self.u.sum(axis=1)[np.newaxis].T, self.n_clusters)

    for iteration in range(self.max_iter):
        u_old = self.u.copy()

        self.centers = self.next_centers(X)
        self.u = self._predict(X)

        selfClusterOut = self.predict(X)
        centers = self.centers

        for i in range(nPoints):
            plt.scatter(x[i], y[i], c = plotColor[selfClusterOut[i]], s = 10)
        for i in range(self.n_clusters):
            plt.scatter(centers[i][0], centers[i][1], c = 'black', marker = 'X')
        plt.show()

        # Stopping rule
        if norm(self.u - u_old) < self.error:
            break

def next_centers(self, X):
    um = self.u ** self.m
    return (X.T @ um / np.sum(um, axis=0)).T

def _predict(self, X):
    power = float(2 / (self.m - 1))
    temp = cdist(X, self.centers) ** power
    denominator_ = temp.reshape((X.shape[0], 1, -1)).repeat(temp.shape[-1], axis=1)
    denominator_ = temp[:, :, np.newaxis] / denominator_

    return 1 / denominator_.sum(2)

def predict(self, X):
    if len(X.shape) == 1:
        X = np.expand_dims(X, axis=0)

    u = self._predict(X)
    return np.argmax(u, axis=-1)

```

In [53]:

```

from google.colab import drive
drive.mount('/content/drive')

```

```

-----
ModuleNotFoundError                                Traceback (most recent call last)
<ipython-input-53-d5df0069828e> in <module>
----> 1 from google.colab import drive
      2 drive.mount('/content/drive')

```

ModuleNotFoundError: No module named 'google'

In [54]:

```
X, y = make_circles(n_samples=750, factor=0.3, noise=0.1)
X = StandardScaler().fit_transform(X)
y_pred = DBSCAN(eps=0.3, min_samples=10).fit_predict(X)
y_kmeans = KMeans(2).fit_predict(X)

plt.scatter(X[:,0], X[:,1], c=y_pred)
plt.show()
plt.scatter(X[:,0], X[:,1], c=y_kmeans)
plt.show()
print('Number of clusters: {}'.format(len(set(y_pred[np.where(y_pred != -1)]))))
print('Homogeneity: {}'.format(metrics.homogeneity_score(y, y_pred)))
print('Completeness: {}'.format(metrics.completeness_score(y, y_pred)))
```

NameError Traceback (most recent call last)

```
<ipython-input-54-c6e26e7f8c07> in <module>
----> 1 X, y = make_circles(n_samples=750, factor=0.3, noise=0.1)
      2 X = StandardScaler().fit_transform(X)
      3 y_pred = DBSCAN(eps=0.3, min_samples=10).fit_predict(X)
      4 y_kmeans = KMeans(2).fit_predict(X)
      5
```

NameError: name 'make_circles' is not defined

In []: